

EvoTrees.jl

Efficient Boosted Trees on
CPUs & GPUs



Jérémie Desgagné-Bouchard, FCAS
Head of Science - Evovest



Motivation

- **Key operations of the Boosted Trees learning algo**
 - And Performance tricks used by EvoTrees.jl
- **Latest features**
 - Tree bagging
 - Oblivious trees
- **Alternatives to gradient-based losses**
 - mean-absolute-error
 - volatility-adjusted losses
- **Future development paths**
 - Improved GPU acceleration
 - Multi-target losses
 - Enhanced support of categorical variables
 - Autodiff support

Basics

MLJ compliant API

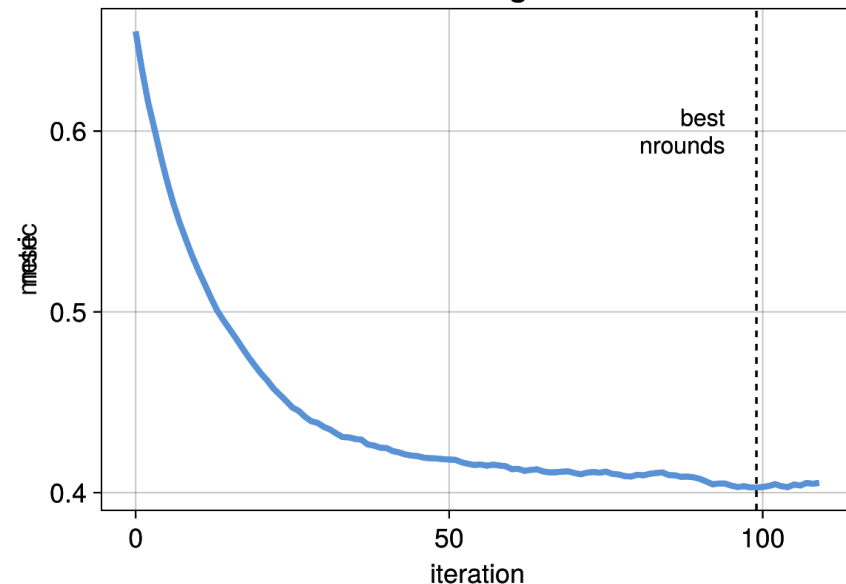
Row	target_name		feature_names						
	Survived Int64	Pclass Int64	Sex Cat...	Age Float64	SibSp Int64	Parch Int64	Fare Float64	Age_ismissing Bool	
1	0	1	male	29.0	0	0	30.0	false	

```
config = EvoTreeRegressor(  
    loss=:logloss,  
    nrounds=200,  
    early_stopping_rounds=10,  
    eta=0.05,  
    max_depth=5,  
)
```

```
model = EvoTrees.fit(  
    config, dtrain;  
    deval,  
    target_name,  
    feature_names,  
    print_every_n=10  
)
```

```
pred = model(deval)
```

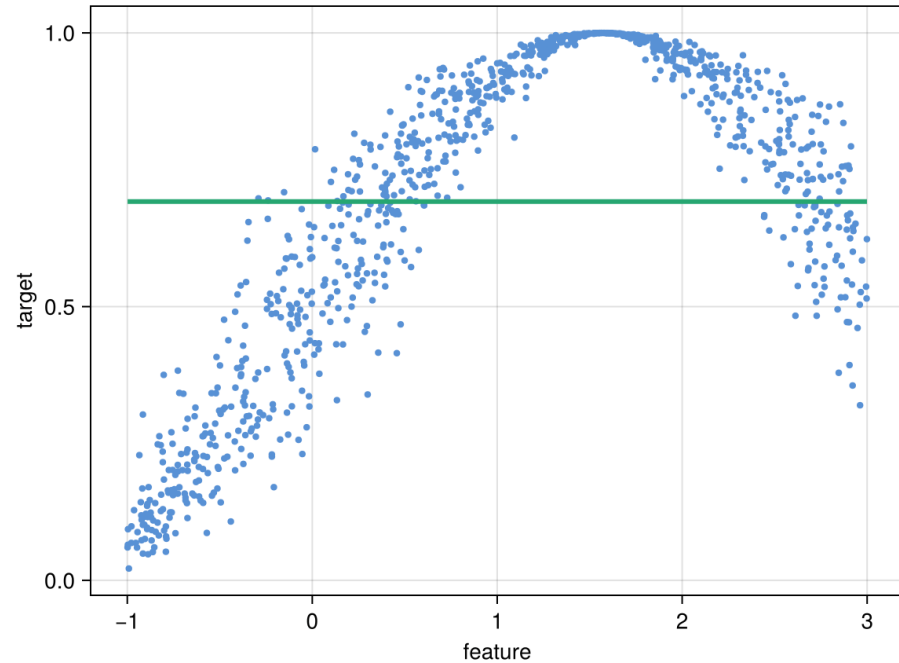
Eval metric through iterations



Algo Overview

Boosted algorithm

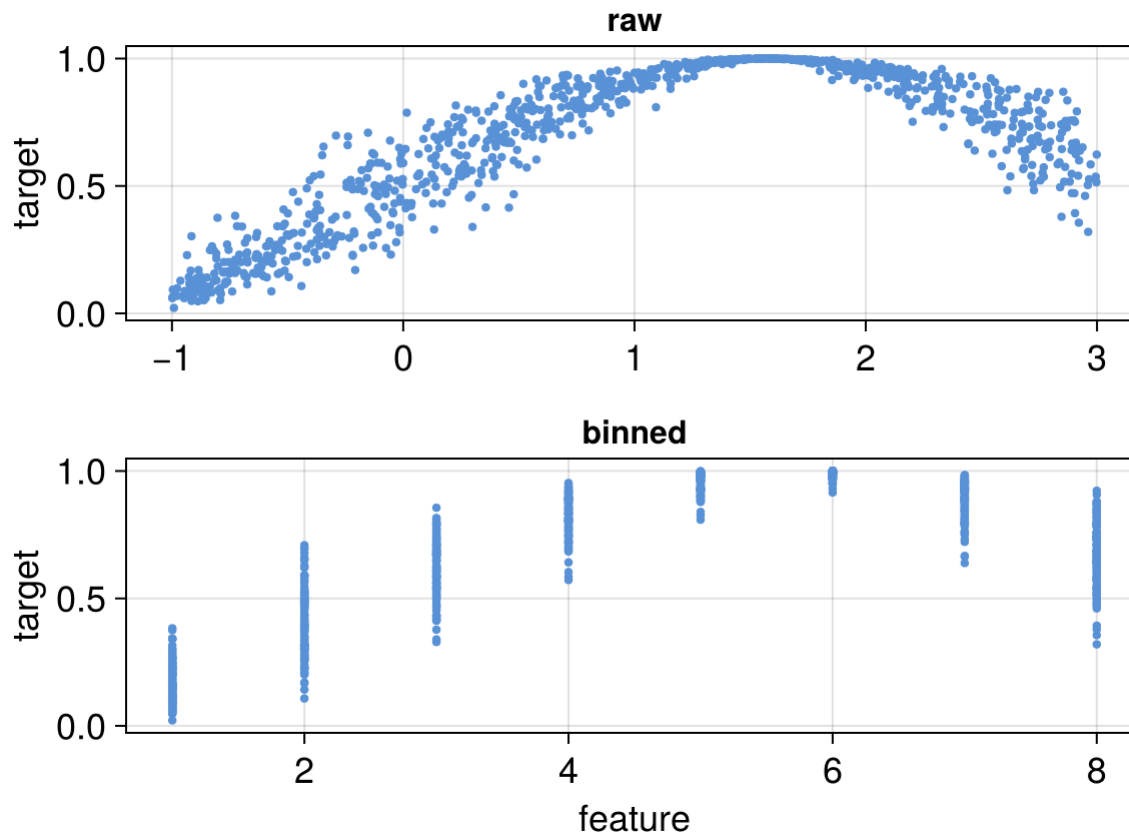
0.692



tree #1 = bias

Algo – Histogram-based

A core methodology for fast training used by all performant implementations (XGBoost, LightGBM, CatBoost)



Note how the binarization results in features that are scale and distribution invariant

Training Speed

CPU

# obs	depth	EvoTrees	XGBoost	log-diff
100K	6	0.83	1.31	-20%
100K	11	3.43	3.33	1%
1M	6	5.50	13.57	-39%
1M	11	15.72	18.79	-8%
10M	6	79.99	147.04	-26%
10M	11	185.47	189.04	-1%

GPU

# obs	depth	EvoTrees	XGBoost	log-diff
100K	6	1.64	0.66	40%
100K	11	17.98	3.77	68%
1M	6	3.53	3.19	4%
1M	11	30.11	8.58	55%
10M	6	21.56	27.49	-11%
10M	11	65.96	49.85	12%

Tree Bagging

Tree bagging allows to train models similar to RandomForest

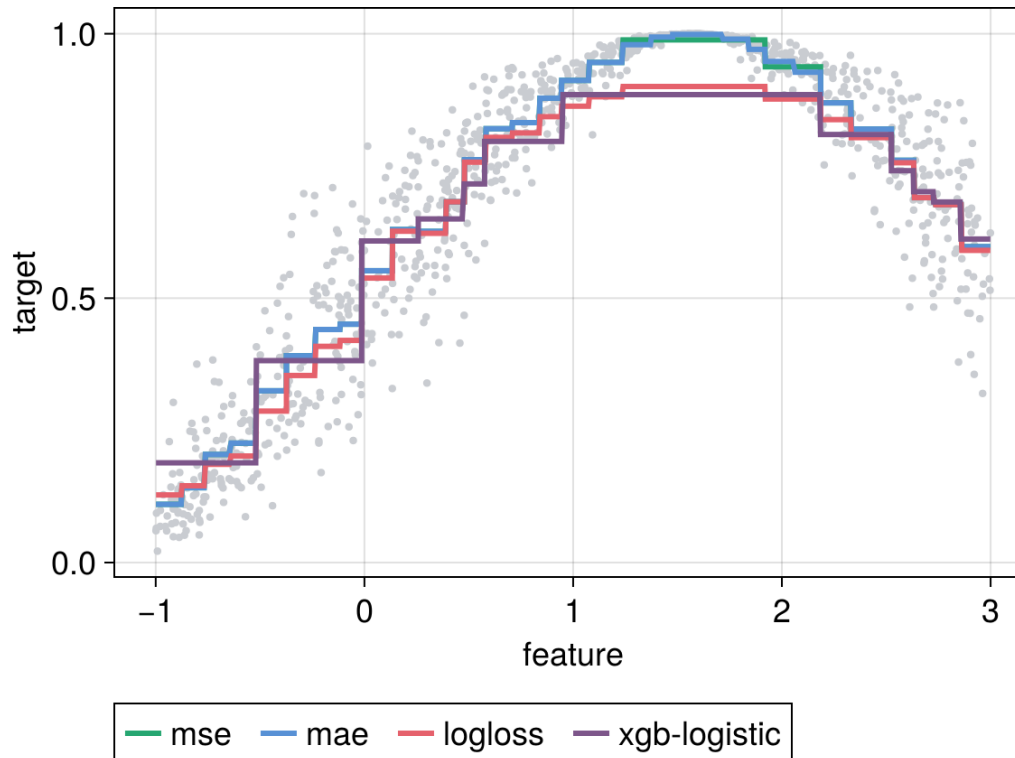
```
for tree in tree_size
  update_gradients!
  for bag in bag_size
    subsample_row_cols!
    grow_tree!
    predict!
  end
end
```

```
EvoTreeRegressor(
  nrounds=1,
  bagging_size=8,
  eta=1.0,
  row_sample=0.5,
  col_sample=0.5,
  loss=:mse,
  kwargs...
)
```

Tree Bagging

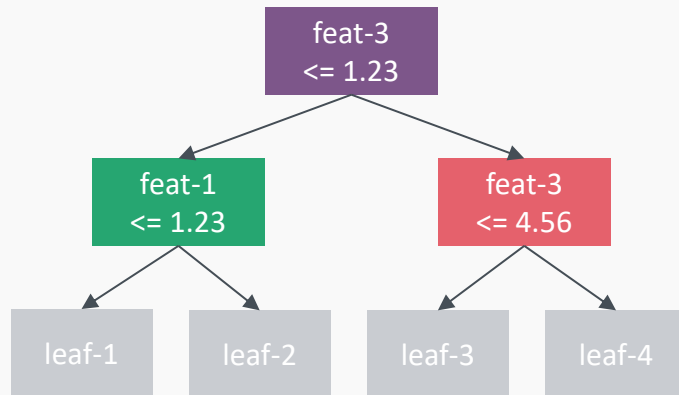
⚠ Only few loss functions can return a good model with a RandomForest emulation: $nrounds = 1$; $eta = 1.0$

Limitation intrinsic to gradient learning method. A single tree quality is limited to the quality of second-order approximation vs. the original loss curve.




Oblivious Tree Structure

Binary



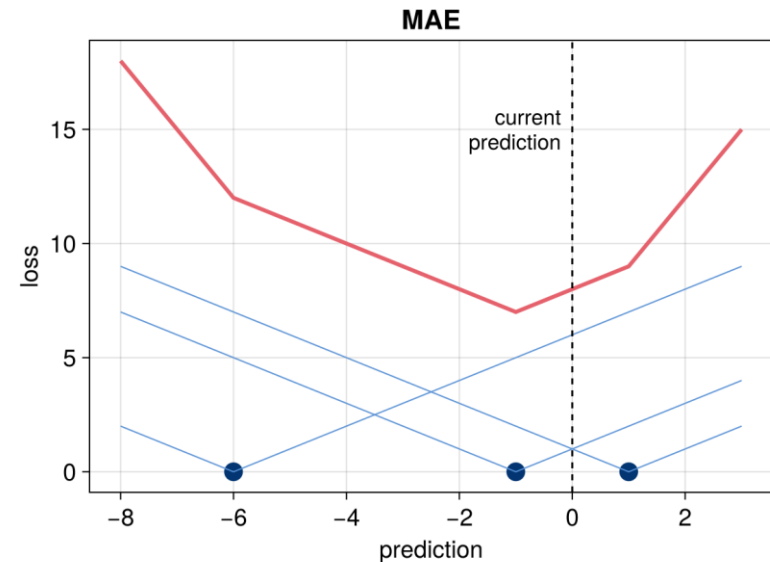
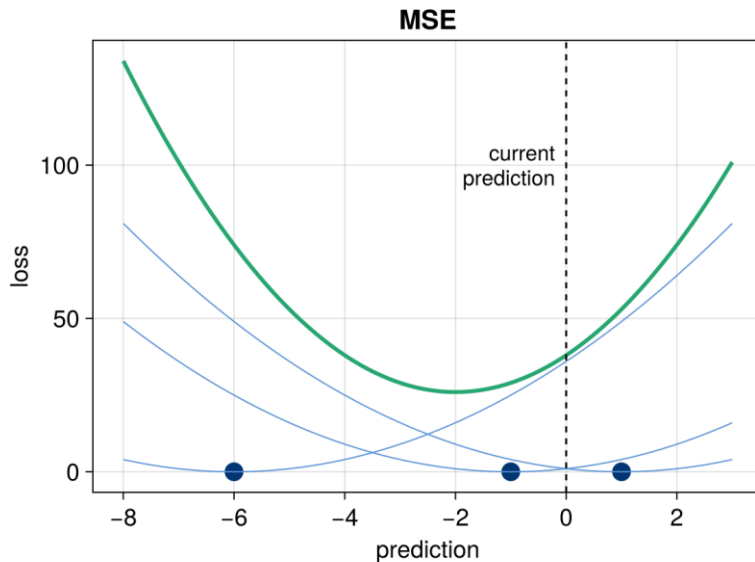
```
EvoTreeRegressor(  
  tree_type = :binary,  
  kwargs...  
)
```

Oblivious



```
EvoTreeRegressor(  
  tree_type = :oblivious,  
  kwargs...  
)
```

Loss Without Gradients - MAE



Are the 2 first moments needed if all we want is to move prediction towards the mean?

Alternative perspective for tree-split gain: any metric derived from histogram statistics.
For *MAE proxy*: gain = sum of absolute change in predicted value.




A single metric is needed, the residual: ***target - pred***

Cannot work beyond depth 2 without a trick gain valuation: gains need to be derived from the context of the parent node statistics.

Loss Without Gradients - Credibility

All other things being equal, the gain for two split candidates will be the same regardless of the volatility of observations within each leaf.

In regular regression, the gain is entirely driven by the mean of observations.



Loss Without Gradients - Credibility

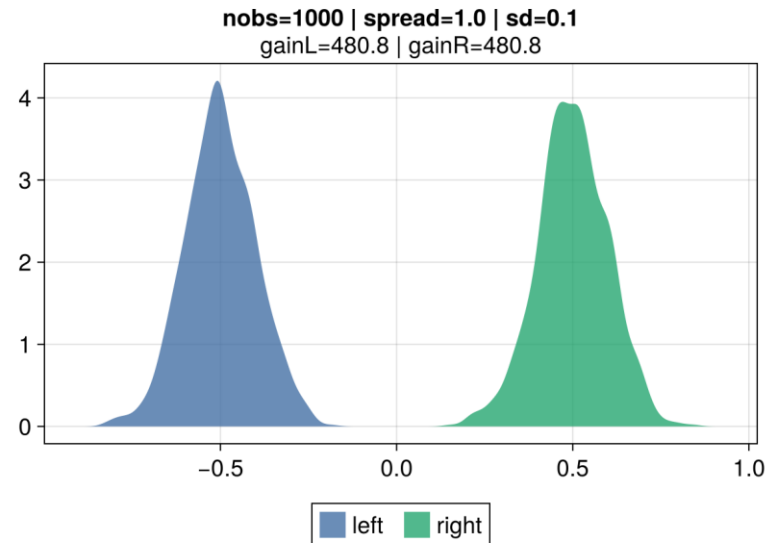

The gain is derived as the credibility-adjusted leaf mean.

Credibility is inversely proportional to the volatility of the leaf's observations.

$$Z = VHM / (VHM + EVPV)$$

Variance of the Hypothetical Means (VHM): if large differences between candidate means are expected, a greater credibility is assigned.

Expected Value of the Process Variance (EVPV): if the data generation process of a given candidate has high volatility, a lower credibility is assigned.



Performance Benchmarks

MSE metrics on regression problems

model	mse	mae	cred_var	cred_std
boston	6.3	5.71	5.95	5.43
year	74.9	74.3	74.6	74.2
msrank	0.55	0.549	0.551	0.549
yahoo	0.565	0.567	0.589	0.568

Roadmap

- **GPU improvement**

- SciML's small grant to bring GPU at XGBoost level (USD 1,500)
- <https://github.com/Evovest/EvoTrees.jl/issues/288>

- **Multi-target support**

- Tree structure where leaves generate a vector of predictions, one for each of the multiple target variables

```
EvoTrees.fit(  
    config,  
    dtrain;  
    target_name = [y1, y2, y3],  
    kwargs...  
)
```

Questions?

